

Web-Scale KMean clustering: Paper Review

By: Simon Bassey

Paper

Author:

D. Sculley Google, Inc. Pittsburgh.
PA USA

Publisher: 2010

Source: [web-scale kmean clustering](#)

Summary of Core Idea

The Lloyd's classic batch k-means clustering algorithm has a computation complexity of $O(kns)$

where

k =number of clusters, n = size of dataset, and

s =max. Number of non-zero elements in any one example in n

And thus, does not scale for large datasets, to meet the near real-time latency requirement of under **1000 milliseconds** for most web applications.



Proposed Experiment

The paper proposed and implemented two approaches to faster clustering of large dataset:

1. **Mini-batch K-Means** clustering algorithm

For the task of predicting clusters for new new examples, with a lower computation complexity cost of

$$\sim O(Kbs) < O(kns),$$

Where $b < n$ is a batch size, $10\times$ in order of magnitude less than n .



Proposed Experiment ..

2. Sparse Cluster Centers

That further achieves better server resource utilization with compact storage of the model and reduces network costs, using an efficient projection-gradient method, which projects the centroids after update, to the nearest point within an L1-ball of some given radius λ .



Implementation Overview

Mini-batch K-Means

```
def train(self, Xs, bs):
    rand_indices = [random.randrange(len(Xs)) for i in range(self.__k)]
    centroids = Xs[rand_indices]

    clusters = [[] for i in range(self.__k)]
    prev_centroids = None
    iter_count = 0;
    while iter_count < self.__m_iter:
        iter_count +=1
        n_clusters = [[] for i in range(self.__k)]
        batch_Xs = Xs[np.random.choice(Xs.shape[0], bs, replace=True)]
        idx_x_c = np.empty(batch_Xs.shape[0], dtype=int)
        V = np.zeros(centroids.shape[0])

        for i,xi in enumerate(batch_Xs):
            eucl_dst = [np.linalg.norm(xi - centroid) for centroid in centroids]
            min_cent_idx = np.argmin(eucl_dst)
            idx_x_c[i] = min_cent_idx
        prev_centroids = centroids
        print(idx_x_c)
        for j, x in enumerate(batch_Xs):
            V[idx_x_c[j]] +=1
            lr = 1.0/V[idx_x_c[j]]
            centroids[idx_x_c[j]] = (1.0 - lr) * centroids[idx_x_c[j]] + lr*x
            n_clusters[idx_x_c[j]].append(xi.tolist())
        optimised = True
        for c in range(self.__k):
            if np.sum((centroids[c]-prev_centroids[c])/prev_centroids[c] *100) > 0.001:
                optimised = False
        if optimised:
            break
    self.centroids = centroids
    return iter_count,centroids,n_clusters,
```

Train with sparse cluster update

<Not Implemented>

Experiment Results

The resulting mini-batch implementation trained on the RCV1 collection of documents dataset of >800K examples, and tested on >23k test data points achieved:

1. **Great Accuracy -**


An accuracy slightly lower than the classic Lloyd's batch k-means clustering algorithm, but beat the SGD and triangular inequality variants of the k-means clustering algorithm.

2. **Lower computational cost -**

Very importantly, the implementation had lower computational cost in comparison as seen in the next slide, suitable for most web Applications.

Suggested improvements

As the paper is relatively old, a number of faster implementations have been proposed over the years, even then new implementations looking to use the recommendations for m the paper for fast k-means may need to keep in that:

1. Faster initialization algorithms for choosing the initial centroids for clusters exist today and could be explored.
 2. For very large dataset on multi-core machines, the task of choosing the closest centroid for all examples in a batch as well as updates could be parallelized, thus reducing further the complexity to $\sim O(ks)$ instead of $O(kbs)$ as proposed.
 3. Although not mentioned in the paper, early stopping could also be explored to stop centroid updates no longer improve whether or not the number of epoch iterations is reached.
- 

Thank you for listening



Questions?